

Importing Data

Daniel Anderson

Agenda

- Discuss reading in data
 - The *rio* package, and when you may want some additional flexibility
 - Other packages for reading in data
- Lab
 - Use R Markdown
 - Read in data from two different sources
 - Conduct some basic manipulations

rio

- Super nice package - most of the time, it just works, regardless of the source file type.
- (If this isn't astounding to you, you obviously haven't struggled for hours to read in data properly)

Example: these all work! Try it out and verify for yourself!

```
library(rio)
exam1 <- import("./data/exam1.csv")
eclsk <- import("./data/ecls-k_samp.sav")
fatality <- import("./data/Fatality.txt")
```

Read in text files directly from the web!

```
reads <- import("https://data.jacksonms.gov/api/views/97iy-g8hk/rows.csv")
head(reads)
```

```
##           Test Year Test Type Test Site Student ID Pre-Test Score
## 1 06/01/2016 12:00:00 AM  YEAR END  VIRDEN  Virден 1           43
## 2 06/01/2016 12:00:00 AM  YEAR END  VIRDEN  Virден 2           46
## 3 06/01/2016 12:00:00 AM  YEAR END  VIRDEN  Virден 3           39
## 4 06/01/2016 12:00:00 AM  YEAR END  VIRDEN  Virден 4           35
## 5 06/01/2016 12:00:00 AM  YEAR END  VIRDEN  Virден 5           46
## 6 06/01/2016 12:00:00 AM  YEAR END  VIRDEN  Virден 6           35
##  Pre-Test (%) Post-Test Score Post-Test (%) Percentage Change
## 1           29%           92           61%           32%
## 2           31%          104           69%           38%
## 3           26%           75           50%           24%
## 4           23%          115           77%           54%
## 5           31%           85           57%           26%
## 6           23%           91           61%           38%
##  Unit 1 Score Unit 1 (%) Unit 2 Score Unit 2 (%) Unit 3 Score Unit 3 (%)
## 1           3          12%           4          16%           6          24%
## 2           5          20%           5          20%           6          24%
## 3           4          16%           4          16%           6          24%
```

You can export just as easily!

Try it out!

```
library(janitor)
reads <- clean_names(reads)
export(reads, "project_reads.sav")
export(reads, "project_reads.txt")
export(reads, "project_reads.dta")
```

Note. The `clean_names` function was necessary because spaces aren't valid for SPSS or Stata variable names. If you don't run `clean_names()` first the stata export will fail, while Running `clean_names` first fixes the issue.

Export any data frame in any format

```
library(tidyverse)
mtcars %>%
  group_by(cyl) %>%
  summarize(mean_mpg_by_cyl = mean(mpg))
```

```
## # A tibble: 3 x 2
##   cyl mean_mpg_by_cyl
##   <dbl>           <dbl>
## 1     4         26.66364
## 2     6         19.74286
## 3     8         15.10000
```

```
mtcars %>%
  group_by(cyl) %>%
  summarize(mean_mpg_by_cyl = mean(mpg)) %>%
  export("mpg_mean_by_cyl.sav")
```

convert ()

- Another really useful feature is `convert ()`, which just takes a file of one type and converts it to another.
- Say your advisor uses SPSS, but her/his colleague uses Stata. They might use some proprietary (and expensive) software like SAS/CONNECT. Instead, just run this one line of code and voila!

```
convert("./data/ec1s-k_samp.sav", "./data/ec1s-k_samp.dta")
```

How is this all working?

Looking at the import documentation

```
?import
```

```
import {rio}
```

R Documentation

Import

Description

Read in a `data.frame` from a file

Usage

```
import(file, format, setclass, which, ...)
```

Arguments

- | | |
|-----------------|--|
| file | A character string naming a file, URL, or single-file .zip or .tar archive. |
| format | An optional character string code of file format, which can be used to override the format inferred from <code>file</code> . Shortcuts include: <code>","</code> (for comma-separated values), <code>";"</code> (for semicolon-separated values), and <code>" "</code> (for pipe-separated values). |
| setclass | An optional character vector specifying one or more classes to set on the import. By default, all the return object is always a <code>"data.frame"</code> . Allowed values for this might be <code>"tbl_df"</code> , <code>"tbl"</code> , or <code>"tibble"</code> (if using <code>dplyr</code>) or <code>"data.table"</code> (if using <code>data.table</code>). Other values are ignored such that a <code>data.frame</code> is returned. |
| which | This argument is used to control import from multi-object files; as a rule <code>import</code> only ever returns a single data frame. (Use import_list to import multiple data frames from a multi-object file.) If <code>file</code> is a compressed directory, <code>which</code> can be either a character string specifying a filename or an integer specifying which file (in locale sort order) to extract from the compressed directory. For Excel spreadsheets, this can be used to specify a sheet number. For .Rdata files, this can be an object name. For HTML files, which table to extract (from document order). Ignored otherwise. A character string value will be used as a regular expression, such that the extracted file is the first match of the regular expression against the file names in the archive. |
| ... | Additional arguments passed to the underlying import functions. For example, this can control column classes for delimited file types, or control the use of <code>haven</code> for Stata and SPSS or <code>readxl</code> for Excel (.xlsx) format. See details below. |

So... let's go look at the original packages more!

Tidyverse packages

- **readr**: Designed for quick and efficient reading/writing of plain text files (csv, tsv, txt, etc)
 - not used by *rio*, but if you're having any trouble with csv's, this is the method I'd recommend moving toward.
- **haven**: Designed to read/write files from SPSS, SAS, and Stata files
 - Used by *rio* but with some differences in how the data are actually read in.

readr

- Primary function is `read_csv`
- Used equivalently to `rio::import`, but only works for csv files
- Note the messages produced, below

```
library(tidyverse)
exam1 <- read_csv("./data/exam1.csv")
```

```
## Parsed with column specification:
## cols(
##   .default = col_integer(),
##   stu_name = col_character(),
##   gender = col_character()
## )
```

```
## See spec(...) for full column specifications.
```

Use it to read in text from a string

```
tmp <- read_csv("Jane, Mary, Bob  
                1, 5, 8  
                4, 2, 6  
                3, 5, 1")  
  
tmp
```

```
## # A tibble: 3 x 3  
##   Jane  Mary  Bob  
##   <int> <int> <int>  
## 1     1     5     8  
## 2     4     2     6  
## 3     3     5     1
```

Skipping lines

(also works with `rio::import`)

If there's notes or blank lines to begin with, you can skip over them.

```
read_csv("Here's a line of garbage  
Here's another with some note that you can see in excel but not here  
That's a silly way to store data  
Next line has the actual data.  
Jane, Mary, Bob  
1, 5, 8  
4, 2, 6  
3, 5, 1",  
        skip = 4)
```

```
## # A tibble: 3 x 3  
##   Jane  Mary  Bob  
##   <int> <int> <int>  
## 1     1     5     8  
## 2     4     2     6  
## 3     3     5     1
```

Column names

If there are no column names, they can be supplied with *col_names*.

```
read_csv("1, 5, 8  
         4, 2, 6  
         3, 5, 1",  
         col_names = c("Jane", "Mary", "Bob"))
```

```
## # A tibble: 3 x 3  
##   Jane  Mary  Bob  
##   <int> <int> <int>  
## 1     1     5     8  
## 2     4     2     6  
## 3     3     5     1
```

Missing values

(also works with `rio::import`) Specify your own missing values.

```
read_csv("1, 5, 8
         4, 2, 999
         999, 5, 1",
         col_names = c("Jane", "Mary", "Bob"),
         na = "999")
```

```
## # A tibble: 3 x 3
##   Jane  Mary  Bob
##   <int> <int> <int>
## 1     1     5     8
## 2     4     2    NA
## 3    NA     5     1
```

Other separators

What if the data are separated by tabs or something like "|"?

- Use `read_delim()` and specify the delimiter.
- Alternatively, specify the format as "|" with `rio`. Ex: `import(file.txt, format = "|")`

```
read_delim("1|5|8
           4|2|999
           999|5|1",
           delim = "|",
           col_names = c("Jane", "Mary", "Bob"),
           na = "999")
```

```
## # A tibble: 3 x 3
##       Jane  Mary  Bob
##   <chr> <int> <int>
## 1      1      5      8
## 2      4      2     NA
## 3    999      5      1
```


Specify column type

By default, these are all read in as integer. Let's change it.

```
read_delim("1|5|8
           4|2|999
           999|5|1",
           delim = "|",
           col_names = c("Jane", "Mary", "Bob"),
           na = "999",
           col_types = cols(
             col_character(),
             col_double(),
             col_integer()))
```

```
## # A tibble: 3 x 3
##       Jane  Mary  Bob
##   <chr> <dbl> <int>
## 1      1      5      8
## 2      4      2     NA
## 3    999      5      1
```

Important things to think about

- When importing data, how many rows and columns do you expect? See if it matches using `dim(ob)` where ob is the data object.
- Be careful of missing data (how are they coded in the original file?)
- Always do some double checking to make sure everything read in correctly
 - e.g., `head()` `tail()`, `summary()`, `str()`

Importing data from other sources

The *haven* package

- Really powerful package - much better than the *foreign* package that comes pre-installed.
- Can read *and write* SPSS, SAS, and Stata files.
- By default, user-defined missing data will be read in as missing.
- Used by *rio* so all arguments should be passed directly



Example

Load a sample of the ECLS-K dataset

```
library(haven) # part of tidyverse so should already be installed
eclsk_haven <- read_sav("./data/ecls-k_samp.sav")
eclsk_haven
```

```
## # A tibble: 984 x 33
##   child_id teacher_id school_id k_type school_type sex ethnic
##   <chr>      <chr>      <chr> <dbl+lbl> <dbl+lbl> <dbl+lbl> <dbl+lbl>
## 1 0842021C 0842T02      0842      1          0          0          2
## 2 0905002C 0905T01      0905      1          1          0          5
## 3 0150012C 0150T01      0150      1          1          1          2
## 4 0556009C 0556T01      0556      1          1          1          4
## 5 0089013C 0089T04      0089      1          0          0          1
## 6 1217001C 1217T13      1217      0          0          1          6
## 7 1092008C 1092T01      1092      0          0          1          4
## 8 0083007C 0083T16      0083      1          0          0          1
## 9 1091005C 1091T02      1091      0          1          0          1
## 10 2006006C 2006T01      2006      1          1          0          1
## # ... with 974 more rows, and 26 more variables: famtype <dbl+lbl>,
## #   numsibs <dbl>, SES_cont <dbl>, SES_cat <dbl+lbl>, age <dbl>,
## #   T1RSCALE <dbl>, T1MSCALE <dbl>, T1GSCALE <dbl>, T2RSCALE <dbl>,
```

labels

- haven tries to maintain the attributes a variable had when it was in SPSS, SAS, or STATA.
- To do this, it provides a new *labelled* class
- This way, no information is lost, and you can decide what to do with them
 - remove labels
 - coerce to factor
- This is slightly different than how `rio::haven` handles the data on import

labelled class

```
##          haven      rio
## child_id  character character
## teacher_id character character
## school_id character character
## k_type    labelled   numeric
## school_type labelled   numeric
## sex       labelled   numeric
## ethnic    labelled   numeric
## famtype   labelled   numeric
## numsibs   numeric    numeric
## SES_cont   numeric    numeric
## SES_cat    labelled   numeric
## age        numeric    numeric
## T1RSCALE   numeric    numeric
## T1MSCALE   numeric    numeric
## T1GSCALE   numeric    numeric
## T2RSCALE   numeric    numeric
## T2MSCALE   numeric    numeric
## T2GSCALE   numeric    numeric
## IRTreadgain numeric    numeric
## IRTmathgain numeric    numeric
## IRTgkgain  numeric    numeric
```

Make them what you want

The main difference between the packages:

(note - these all work on a full data frame as well as individual columns)

If you want a variable to be numeric

- rio - nothing
- `haven::zap_labels()`

If you want a variable to be a factor (we'll talk more about these later)

- `rio::factorize()`
- `haven::as_factor()`

If you want a variable to be a character (more complicated with haven)

- `rio::characterize()`

Final notes for importing data

- Generally reading in data is not a big deal. Occasionally tricky formats can come up.
- `rio::import` should work most of the time - particularly in this class - but you may want need to pass additional arguments at times.
 - One of the drawbacks (but makes things simpler) is that you may not know which variables had labels attached and which did not.