

Factors and Dates

Daniel Anderson

Agenda

- Basics of factors
 - creating/modifying
 - when you do/do not want factors
- Basics of Dates
 - Specifically, we'll focus date calculations

Disclaimer

- We're obviously not going to cover all there is to know about factors and dates in one smashed-together two-hour lecture.
- If we had more time, we'd spend a week on each. Instead you get one lecture.

Factors

Notice a difference?

```
library(tidyverse)
tibble(lets = letters[1:3])
```

```
## # A tibble: 3 x 1
##   lets
##   <chr>
## 1     a
## 2     b
## 3     c
```

```
data.frame(lets = letters[1:3])
```

```
##   lets
## 1     a
## 2     b
## 3     c
```

What about now?

```
str(tibble(lets = letters[1:3]))
```

```
## Classes 'tbl_df', 'tbl' and 'data.frame':    3 obs. of  1 variable:  
## $ lets: chr  "a" "b" "c"
```

```
str(data.frame(lets = letters[1:3]))
```

```
## 'data.frame':    3 obs. of  1 variable:  
## $ lets: Factor w/ 3 levels "a","b","c": 1 2 3
```

Why?

- Primarily historical reasons
 - Factors used to be much easier to work with
 - If you want to use the data for modeling, factors make more sense
 - R is increasingly used for all sorts of things besides analysis, so it makes less sense for everything to be a factor

What to do?

- Turn it off globally, but that's dangerous

```
options(default.stringsAsFactors = FALSE)
```

- Turn it off in only the functions it affects, but you might forget

```
str(data.frame(lets = letters[1:3], stringsAsFactors = FALSE))
```

```
## 'data.frame':    3 obs. of  1 variable:  
## $ lets: chr  "a" "b" "c"
```

- Use `rio::import` or `readr` (e.g., `readr::read_csv`), which will default to reading strings in as strings

Creating factors

- Imagine you have a vector of months

```
months <- c("Dec", "Apr", "Jan", "Mar")
```

- We could store this as a string, but there are issues with this.
 - There are only 12 possible months
 - factors will help us weed out values that don't conform to our predefined *levels*, which helps safeguard against typos, etc.
 - You can't sort this vector in a meaningful way (it defaults to alphabetic)

```
sort(months)
```

```
## [1] "Apr" "Dec" "Jan" "Mar"
```

Define it as a factor

```
month_levels <- c("Jan", "Feb", "Mar", "Apr", "May", "Jun",  
                  "Jul", "Aug", "Sep", "Oct", "Nov", "Dec")
```

```
months <- factor(months, levels = month_levels)  
months
```

```
## [1] Dec Apr Jan Mar  
## Levels: Jan Feb Mar Apr May Jun Jul Aug Sep Oct Nov Dec
```

- Now, we can sort

```
sort(months)
```

```
## [1] Jan Mar Apr Dec  
## Levels: Jan Feb Mar Apr May Jun Jul Aug Sep Oct Nov Dec
```

Also provides a safety net of sorts

```
months[5] <- "Jam"
```

```
## Warning in `[<-.factor`(`*tmp*`, 5, value = "Jam"): invalid factor level,  
## NA generated
```

```
months
```

```
## [1] Dec  Apr  Jan  Mar  <NA>  
## Levels: Jan Feb Mar Apr May Jun Jul Aug Sep Oct Nov Dec
```

What if we don't specify the levels?

- If you define a factor without specifying the levels, it will assign them alphabetically

```
mnths <- c("Dec", "Apr", "Jan", "Mar")  
factor(mnths)
```

```
## [1] Dec Apr Jan Mar  
## Levels: Apr Dec Jan Mar
```

- If you instead want them in the order they appeared in the data, use **unique** when specifying the levels (Why is **unique()** necessary? What's it doing?)

```
factor(mnths, levels = unique(mnths))
```

```
## [1] Dec Apr Jan Mar  
## Levels: Dec Apr Jan Mar
```

Accessing and modifying levels

Use the **levels** function

- To view the levels

```
levels(months)
```

```
## [1] "Jan" "Feb" "Mar" "Apr" "May" "Jun" "Jul" "Aug" "Sep" "Oct" "Nov"  
## [12] "Dec"
```

- To modify the levels

```
levels(months) <- 1:12  
months
```

```
## [1] 12  4  1  3  <NA>  
## Levels: 1 2 3 4 5 6 7 8 9 10 11 12
```

If you need to, be specific

```
months <- factor(months, levels = 1:12, labels = month_levels)
months
```

```
## [1] Dec  Apr  Jan  Mar  <NA>
## Levels: Jan Feb Mar Apr May Jun Jul Aug Sep Oct Nov Dec
```

New package

- When working with factors, we can use the *forcats* package
 - *for* categorical variables
 - anagram for factors
- Part of the tidyverse
 - Should be installed for you already, but won't load with `library(tidyverse)`



Changes factors back to the order they appeared

```
c("Dec", "Apr", "Jan", "Mar") %>%  
  factor()
```

```
## [1] Dec Apr Jan Mar  
## Levels: Apr Dec Jan Mar
```

```
c("Dec", "Apr", "Jan", "Mar") %>%  
  factor(levels = c("Jan", "Mar", "Apr", "Dec"))
```

```
## [1] Dec Apr Jan Mar  
## Levels: Jan Mar Apr Dec
```

see next slide


```
library(forcats)
c("Dec", "Apr", "Jan", "Mar") %>%
  factor(levels = c("Jan", "Mar", "Apr", "Dec")) %>%
  fct_inorder()
```

```
## [1] Dec Apr Jan Mar
## Levels: Dec Apr Jan Mar
```

Or order by frequency

```
c("b", "b", "c", "a", "a", "a") %>%  
  fct_infreq()
```

```
## [1] b b c a a a  
## Levels: a b c
```

- This can be particularly useful for plotting


words example

```
data(sentences, package = "stringr")
sentences <- tibble(sent_num = seq_along(sentences), sentence = sentences)
library(tidytext)
words_freq <- sentences %>%
  unnest_tokens(word, sentence) %>%
  count(word) %>%
  filter(n > 30)
words_freq
```

```
## # A tibble: 13 x 2
##   word      n
##   <chr> <int>
## 1     a    202
## 2   and    118
## 3   for     35
## 4    he     34
## 5    in     87
## 6    is     81
## 7    it     36
## 8    of    132
## 9    on     60
```

Try to plot frequencies

```
ggplot(words_freq, aes(word, n)) +  
  geom_col()
```




Reorder according to frequency

```
words_freq2 <- sentences %>%  
  unnest_tokens(word, sentence) %>%  
  mutate(word = fct_infreq(word)) %>%  
  count(word) %>%  
  filter(n > 30)  
words_freq2
```

```
## # A tibble: 13 x 2  
##       word      n  
##   <fctr> <int>  
## 1    the    751  
## 2     a    202  
## 3    of    132  
## 4    to    123  
## 5   and    118  
## 6    in     87  
## 7    is     81  
## 8   was     66  
## 9    on     60  
## 10 with     51  
## 11   it     36
```

Reproduce plot

```
ggplot(words_freq2, aes(word, n)) +  
  geom_col()
```



Looking at the levels

```
levels(factor(words_freq$word))
```

```
## [1] "a"      "and"    "for"    "he"     "in"     "is"     "it"     "of"     "on"     "the"
## [11] "to"     "was"    "with"
```

```
levels(words_freq2$word)
```

```
## [1] "the"      "a"        "of"        "to"        "and"
## [6] "in"       "is"       "was"       "on"        "with"
## [11] "it"       "for"      "he"        "are"       "from"
## [16] "will"     "his"      "we"        "at"        "but"
## [21] "were"     "into"     "they"      "you"       "your"
## [26] "that"     "when"     "this"      "by"        "be"
## [31] "old"      "than"     "as"        "high"      "out"
## [36] "red"      "there"    "these"     "down"      "fine"
## [41] "green"    "hot"      "new"       "she"       "small"
## [46] "strong"   "up"       "used"      "wall"      "before"
## [51] "good"     "hard"     "her"       "makes"     "round"
## [56] "thin"     "two"      "water"     "way"       "young"
## [61] "best"     "blue"     "both"      "bright"    "dull"
## [66] "each"     "gold"     "him"       "kept"      "last"
```

When do we really want factors?

Generally two reasons to declare a factor

- Only finite number of categories
 - Treatment/control
 - Income categories
 - Performance levels
 - etc.
- Use in modeling

GSS

General Social Survey

- We dealt with some of these data for a homework.
- Unbeknownst to me, Hadley also included a sample in the *forcats* dataset

```
gss_cat
```

```
## # A tibble: 21,483 x 9
##   year      marital    age  race      rincome      partyid
##   <int>      <fctr> <int> <fctr>      <fctr>      <fctr>
## 1  2000 Never married    26 White  $8000 to 9999 Ind,near rep
## 2  2000 Divorced        48 White  $8000 to 9999 Not str republican
## 3  2000 Widowed         67 White Not applicable Independent
## 4  2000 Never married    39 White Not applicable Ind,near rep
## 5  2000 Divorced        25 White Not applicable Not str democrat
## 6  2000 Married         25 White $20000 - 24999 Strong democrat
## 7  2000 Never married    36 White $25000 or more Not str republican
## 8  2000 Divorced        44 White  $7000 to 7999 Ind,near dem
## 9  2000 Married         44 White $25000 or more Not str democrat
## 10 2000 Married         47 White $25000 or more Strong republican
## # ... with 21,473 more rows, and 3 more variables: relig <fctr>,
```

Investigate factors

Tidyverse gives you convenient ways to evaluate factors

- Use `count` - no need to use `group_by`
- Use `geom_bar` or `geom_col` with *ggplot*


```
gss_cat %>%
  count(partyid)
```

```
## # A tibble: 10 x 2
##           partyid      n
##           <fctr> <int>
## 1      No answer    154
## 2    Don't know      1
## 3    Other party    393
## 4 Strong republican 2314
## 5 Not str republican 3032
## 6    Ind,near rep   1791
## 7    Independent   4119
## 8    Ind,near dem   2499
## 9    Not str democrat 3690
## 10 Strong democrat 3490
```


```
gss_cat %>%
  count(relig)
```

```
## # A tibble: 15 x 2
##           relig      n
##           <fctr> <int>
## 1      No answer    93
## 2    Don't know    15
## 3 Inter-nondenominational 109
## 4      Native american    23
## 5      Christian    689
## 6 Orthodox-christian    95
## 7      Moslem/islam   104
## 8    Other eastern    32
## 9      Hinduism     71
## 10     Buddhism    147
## 11      Other     224
## 12      None   3523
## 13      Jewish    388
## 14     Catholic   5124
## 15    Protestant 10846
```

```
ggplot(gss_cat, aes(partyid)) +  
  geom_bar()
```




```
ggplot(gss_cat, aes(relig)) +  
  geom_bar()
```




Include missing categories

```
ggplot(gss_cat, aes(race)) +  
  geom_bar()
```




```
ggplot(gss_cat, aes(race)) +  
  geom_bar() +  
  scale_x_discrete(drop = FALSE)
```



What about this?

```
ggplot(gss_cat, aes(rincome)) +  
  geom_bar()
```



```
levels(gss_cat$rincome)
```


```
## [1] "No answer"      "Don't know"    "Refused"       "$25000 or more"  
## [5] "$20000 - 24999" "$15000 - 19999" "$10000 - 14999" "$8000 to 9999"  
## [9] "$7000 to 7999"  "$6000 to 6999" "$5000 to 5999"  "$4000 to 4999"  
## [13] "$3000 to 3999"  "$1000 to 2999" "Lt $1000"      "Not applicable"
```

```
gss <- gss_cat %>%  
  mutate(rincome = factor(rincome, levels = levels(rincome)[c(15:1, 16)]))  
levels(gss$rincome)
```

```
## [1] "Lt $1000"        "$1000 to 2999"  "$3000 to 3999"  "$4000 to 4999"  
## [5] "$5000 to 5999"  "$6000 to 6999"  "$7000 to 7999"  "$8000 to 9999"  
## [9] "$10000 - 14999" "$15000 - 19999" "$20000 - 24999" "$25000 or more"  
## [13] "Refused"        "Don't know"     "No answer"      "Not applicable"
```




```
ggplot(gss, aes(rincome)) +  
  geom_bar() +  
  coord_flip()
```



Quick aside (and somewhat controversial)

Lollypop charts!




code

```
counts <- gss %>%  
  count(rincome)  
ggplot(counts, aes(rincome, n)) +  
  geom_segment(aes(x = rincome, xend = rincome,  
                  y = 0, yend = n),  
              col = "gray80") +  
  geom_point(size = 3, col = "turquoise") +  
  coord_flip() +  
  theme_classic()
```

Reorder factors


The `forcats::fct_reorder` function allows you to easily reorder factors according to another variable

```
relig_summary <- gss_cat %>%  
  group_by(relig) %>%  
  summarise(age = mean(age, na.rm = TRUE),  
            tvhours = mean(tvhours, na.rm = TRUE),  
            n = n())  
  
ggplot(relig_summary, aes(tvhours, relig)) + geom_point()
```



Note - you could actually include the factor reorder right within the `ggplot` call.


```
relig_summary <- relig_summary %>%  
  mutate(relig = fct_reorder(relig, tvhours))  
  
ggplot(relig_summary, aes(tvhours, relig)) + geom_point()
```



Revisiting our word frequency example

- An easier way to do what we did before, would be to just include the reorder call right within the call to ggplot

```
ggplot(words_freq, aes(fct_reorder(word, n), n)) +  
  geom_col()
```



More on modifying factor levels

- The `forcats::fct_recode` function can make modifying factors more explicit

```
gss_cat %>%  
  mutate(partyid = fct_recode(partyid,  
    "Republican, strong"    = "Strong republican",  
    "Republican, weak"     = "Not str republican",  
    "Independent, near rep" = "Ind,near rep",  
    "Independent, near dem" = "Ind,near dem",  
    "Democrat, weak"       = "Not str democrat",  
    "Democrat, strong"     = "Strong democrat")) %>%  
  count(partyid)
```

```
## # A tibble: 10 x 2  
##           partyid      n  
##           <fctr> <int>  
## 1      No answer   154  
## 2    Don't know     1  
## 3    Other party   393  
## 4 Republican, strong 2314  
## 5 Republican, weak  3032  
## 6 Independent, near rep 1791
```

But this can be pretty easily done with base code too

```
levels(gss_cat$partyid)
```

```
## [1] "No answer"          "Don't know"          "Other party"
## [4] "Strong republican"   "Not str republican"   "Ind,near rep"
## [7] "Independent"         "Ind,near dem"         "Not str democrat"
## [10] "Strong democrat"
```

```
levels(gss_cat$partyid)[c(4:6, 8:10)] <- c("Republican, strong",
      "Republican, weak", "Independent, near rep", "Independent, near dem",
      "Democrat, weak", "Democrat, strong")
levels(gss_cat$partyid)
```

```
## [1] "No answer"          "Don't know"
## [3] "Other party"         "Republican, strong"
## [5] "Republican, weak"     "Independent, near rep"
## [7] "Independent"         "Independent, near dem"
## [9] "Democrat, weak"      "Democrat, strong"
```


Collapsing levels

- `fct_recode` can also be used to collapse levels easily

```
gss_cat %>%  
  mutate(partyid = fct_recode(partyid,  
    "Republican, strong"    = "Strong republican",  
    "Republican, weak"     = "Not str republican",  
    "Independent, near rep" = "Ind,near rep",  
    "Independent, near dem" = "Ind,near dem",  
    "Democrat, weak"       = "Not str democrat",  
    "Democrat, strong"     = "Strong democrat",  
    "Other"                = "No answer",  
    "Other"                = "Don't know",  
    "Other"                = "Other party")) %>%  
  count(partyid)
```

```
## # A tibble: 8 x 2
##           partyid      n
##           <fctr> <int>
## 1           Other    548
## 2 Republican, strong 2314
## 3 Republican, weak  3032
## 4 Independent, near rep 1791
## 5           Independent 4119
## 6 Independent, near dem 2499
## 7 Democrat, weak    3690
## 8 Democrat, strong   3490
```

Or with base syntax

```
data(gss_cat)
levels(gss_cat$partyid)
```

```
## [1] "No answer"          "Don't know"          "Other party"
## [4] "Strong republican"   "Not str republican"   "Ind,near rep"
## [7] "Independent"         "Ind,near dem"         "Not str democrat"
## [10] "Strong democrat"
```

```
levels(gss_cat$partyid)[-7] <- c("other", "other", "other",
  "Republican, strong", "Republican, weak",
  "Independent, near rep", "Independent, near dem",
  "Democrat, weak", "Democrat, strong")
```

Collapse a lot of categories

- In my mind, the most useful functions in *forcats* are for collapsing a lot of categories.
- For example, collapse all categories into *republican*, *democrat*, *independent*, or *other*.

```
gss_cat %>%  
  mutate(partyid = fct_collapse(partyid,  
    other = c("No answer", "Don't know", "Other party"),  
    rep = c("Strong republican", "Not str republican"),  
    ind = c("Ind,near rep", "Independent", "Ind,near dem"),  
    dem = c("Not str democrat", "Strong democrat")  
  ) %>%  
  count(partyid)
```

```
## # A tibble: 4 x 2  
##   partyid      n  
##   <fctr> <int>  
## 1   other   548  
## 2    rep  5346  
## 3    ind  8409  
## 4    dem  7180
```

Sometimes even better

- We can "lump" a bunch of categories together using `fct_lump`.
- Default behavior of `fct_lump` is to create an "other" group that includes all the smallest groups while maintaining "other" as the smallest group represented.
- Can also take optional *n* argument, where *n* represents the number of groups to collapse to

```
gss_cat %>%  
  mutate(rel = fct_lump(relig)) %>%  
  count(rel)
```

```
## # A tibble: 2 x 2  
##       rel      n  
##   <fctr> <int>  
## 1 Protestant 10846  
## 2      Other 10637
```

Collapse to 10 religious groups

```
gss_cat %>%  
  mutate(rel = fct_lump(relig, n = 10)) %>%  
  count(rel)
```

```
## # A tibble: 10 x 2  
##           rel      n  
##      <fctr> <int>  
## 1 Inter-nondenominational 109  
## 2           Christian    689  
## 3   Orthodox-christian    95  
## 4       Moslem/islam    104  
## 5           Buddhism    147  
## 6             None  3523  
## 7           Jewish    388  
## 8           Catholic  5124  
## 9       Protestant 10846  
## 10            Other    458
```

One last thing...

Factors with modeling

```
colors <- factor(c("black", "green", "blue", "blue", "black"))
```

- No need for multiple variables to define a categorical variable: internal dummy-coding
- Change the reference group by defining a new contrast matrix. For example, we can set green to the reference group with the following code.

```
contrasts(colors)
```

```
##           blue green
## black      0      0
## blue       1      0
## green      0      1
```

```
contrasts(colors) <- matrix(c(1, 0,
                              0, 1,
                              0, 0),
                             byrow = TRUE,
                             ncol = 2)
```

Contrast coding (continued)

Alternatively, use some of the built in functions for defining new contrasts matrices

```
contr.helmert(3)
```

```
##      [,1] [,2]  
## 1     -1  -1  
## 2      1  -1  
## 3      0   2
```

```
contr.sum(3)
```

```
##      [,1] [,2]  
## 1      1   0  
## 2      0   1  
## 3     -1  -1
```

```
contrasts(colors) <- contr.helmert(3)  
contrasts(colors)
```

```
##           [,1] [,2]  
## black     -1  -1  
## blue      1  -1  
## green      0   2
```

```
contrasts(colors) <- contr.sum(3)  
contrasts(colors)
```

```
##           [,1] [,2]  
## black      1   0  
## blue       0   1  
## green     -1  -1
```

(see:

http://www.ats.ucla.edu/stat/r/library/contrast_coding.htm)

Pause...

Questions?

Dates

Intro

- Dates are hard - harder than they might seem
- Base syntax can be tricky
- Lots of different packages for helping with dates and time-series data
- We'll focus on the tidyverse version: *lubridate*

Three different types of "Dates"

- date
- date-time (POSIXct)
- time (doesn't have its own class, *hms* package can help here, if you need it)

POSIXct data are much more complicated than dates, so use regular dates if possible.

Date variables look like this:

```
library(lubridate)  
today()
```

```
## [1] "2017-05-29"
```

- This is the standard ISO date format: YYYY-MM-DD.
- Any date variable you have, in any format, will end up looking like this after you convert it to a date.

POSIXct/date-time variables look like this:

```
now( )
```

```
## [1] "2017-05-29 12:47:14 PDT"
```

- Notice they include the date, but also the specific time (in military/24 hour format), down to the specific second.
- Also includes the timezone, which is of course important if you're dealing with seconds of data.

Creating dates

- When you read in data, the dates are likely to be in all sorts of different formats.
- Hopefully, they're at least consistent within a column
- *lubridate* makes individual conversions relatively easy.

```
ymd("2012/02/14")
```

```
## [1] "2012-02-14"
```

```
mdy("03/10/2015")
```

```
## [1] "2015-03-10"
```

```
mdy("03 10 15")
```

```
## [1] "2015-03-10"
```

Conversions

`ymd()`

`ydm()`

`mdy()`

`myd()`

`dmy()`

`dym()`

`yq()`

Need to convert a date-time?

```
mdy_hms("04/16/12 11:48.32 AM")
```

```
## [1] "2012-04-16 11:48:32 UTC"
```

Enforce a time zone

```
mdy_hms("04/16/12 11:48.32 AM", tz = "America/Los_Angeles")
```

```
## [1] "2012-04-16 11:48:32 PDT"
```

```
mdy_hms("04/16/12 11:48.32 AM", tz = "America/New_York")
```

```
## [1] "2012-04-16 11:48:32 EDT"
```

Switch between types

```
as_datetime(today())
```

```
## [1] "2017-05-29 UTC"
```

```
as_date(now())
```

```
## [1] "2017-05-29"
```

Numerical dates

- Sometimes you'll run up against dates like **16750** or **-1250**
- These are number deviating from the "Unix Epoch", which is 1970-01-01

```
as_date(4380) # interpreted as days
```

```
## [1] "1981-12-29"
```

```
as_datetime(4380) # interpreted as seconds
```

```
## [1] "1970-01-01 01:13:00 UTC"
```

Parsing other formats

- What format is *dep_time* in?

```
library(nycflights13)
flights
```

```
## # A tibble: 336,776 x 19
##   year month   day dep_time sched_dep_time dep_delay arr_time
##   <int> <int> <int>   <int>         <int>         <dbl>   <int>
##  1  2013     1     1     517           515           2     830
##  2  2013     1     1     533           529           4     850
##  3  2013     1     1     542           540           2     923
##  4  2013     1     1     544           545          -1    1004
##  5  2013     1     1     554           600          -6     812
##  6  2013     1     1     554           558          -4     740
##  7  2013     1     1     555           600          -5     913
##  8  2013     1     1     557           600          -3     709
##  9  2013     1     1     557           600          -3     838
## 10  2013     1     1     558           600          -2     753
## # ... with 336,766 more rows, and 12 more variables: sched_arr_time <int>,
## #   arr_delay <dbl>, carrier <chr>, flight <int>, tailnum <chr>,
## #   origin <chr>, dest <chr>, air_time <dbl>, distance <dbl>, hour <dbl>,
```

So how could we parse this to be usable?

```
unique(flights$dep_time)
```

```
##      [1]  517  533  542  544  554  555  557  558  559  600  601  602  606
##     [14]  607  608  611  613  615  622  623  624  627  628  629  632  635
##     [27]  637  639  643  644  645  646  651  652  653  655  656  657  658
##     [40]  659  701  702  709  711  712  715  717  719  723  724  725  727
##     [53]  728  729  732  733  734  739  741  743  745  746  749  752  753
##     [66]  754  758  759  800  801  803  804  805  807  809  810  811  812
##     [79]  813  814  817  820  821  822  823  824  825  826  828  829  830
##     [92]  831  832  833  835  839  840  846  848  851  852  853  855  856
##    [105]  857  858  859  900  902  903  904  905  906  908  909  912  913
##    [118]  914  917  920  921  923  926  927  929  930  931  932  933  936
##    [131]  937  940  941  946  947  950  953  955  956  957  959 1003 1005
##    [144] 1007 1009 1010 1011 1021 1024 1025 1026 1028 1029 1030 1031 1032
##    [157] 1033 1037 1038 1042 1044 1047 1048 1053 1054 1056 1058 1059 1101
##    [170] 1103 1105 1107 1109 1111 1112 1113 1114 1120 1123 1124 1125 1127
##    [183] 1128 1130 1132 1133 1135 1137 1143 1144 1147 1150 1153 1154 1155
##    [196] 1157 1158 1200 1202 1203 1204 1205 1206 1208 1211 1217 1219 1220
##    [209] 1222 1228 1230 1231 1237 1238 1240 1241 1245 1246 1248 1251 1252
##    [222] 1253 1255 1257 1258 1301 1302 1304 1305 1306 1310 1314 1315 1316
##    [235] 1317 1318 1320 1323 1325 1327 1330 1333 1336 1337 1339 1341 1342
##    [248] 1343 1344 1346 1350 1351 1353 1354 1355 1356 1358 1400 1402 1408
##    [261] 1411 1416 1418 1419 1421 1422 1423 1424 1428 1430 1431 1433 1436
```

The way I'd probably do it

```
flights %>%  
  mutate(dep_time = stringr::str_pad(dep_time, 4, pad = "0")) %>%  
  separate(dep_time, c("dep_hour", "dep_minute"), 2, convert = TRUE)
```

```
## # A tibble: 336,776 x 20  
##   year month   day dep_hour dep_minute sched_dep_time dep_delay arr_time  
##   * <int> <int> <int>   <int>     <int>         <int>     <dbl>   <int>  
## 1  2013     1     1         5         17           515         2     830  
## 2  2013     1     1         5         33           529         4     850  
## 3  2013     1     1         5         42           540         2     923  
## 4  2013     1     1         5         44           545        -1    1004  
## 5  2013     1     1         5         54           600        -6     812  
## 6  2013     1     1         5         54           558        -4     740  
## 7  2013     1     1         5         55           600        -5     913  
## 8  2013     1     1         5         57           600        -3     709  
## 9  2013     1     1         5         57           600        -3     838  
## 10 2013     1     1         5         58           600        -2     753  
## # ... with 336,766 more rows, and 12 more variables: sched_arr_time <int>,  
## #   arr_delay <dbl>, carrier <chr>, flight <int>, tailnum <chr>,  
## #   origin <chr>, dest <chr>, air_time <dbl>, distance <dbl>, hour <dbl>,  
## #   minute <dbl>, time_hour <dtm>
```

How they handle it in the book

Modulo operators

- `%/`: Integer division
- `%`: Remainder

```
123 %/ 100
```

```
## [1] 1
```

```
123 % 100
```

```
## [1] 23
```

```

flights %>%
  mutate(dep_hour = dep_time %/% 100,
         dep_min = dep_time %% 100) %>%
  select(tailnum, dep_time, dep_hour, dep_min)

```

```

## # A tibble: 336,776 x 4
##   tailnum dep_time dep_hour dep_min
##   <chr>    <int>    <dbl>  <dbl>
## 1  N14228      517        5      17
## 2  N24211      533        5      33
## 3  N619AA      542        5      42
## 4  N804JB      544        5      44
## 5  N668DN      554        5      54
## 6  N39463      554        5      54
## 7  N516JB      555        5      55
## 8  N829AS      557        5      57
## 9  N593JB      557        5      57
## 10 N3ALAA      558        5      58
## # ... with 336,766 more rows

```


Creating dates from multiple variables

- Take a minute... How might you think we could create a single date variable?

```
flights %>%  
  select(year, month, day, hour, minute)
```

```
## # A tibble: 336,776 x 5  
##   year month   day hour minute  
##   <int> <int> <int> <dbl> <dbl>  
## 1  2013     1     1     5     15  
## 2  2013     1     1     5     29  
## 3  2013     1     1     5     40  
## 4  2013     1     1     5     45  
## 5  2013     1     1     6      0  
## 6  2013     1     1     5     58  
## 7  2013     1     1     6      0  
## 8  2013     1     1     6      0  
## 9  2013     1     1     6      0  
## 10 2013     1     1     6      0  
## # ... with 336,766 more rows
```

Nice *lubridate* functions

- `make_date()` and `make_datetime()` functions that can save us a boatload of time.
- Arguments are: year, month, day, hour, min, sec, and tz.
- All arguments have defaults, which are: `1970L`, `1L`, `0L`, `0L`, `0`, and `"UTC"`

```
flights %>%  
  select(year, month, day, hour, minute) %>%  
  mutate(departure = make_datetime(year, month, day, hour, minute))
```

```
## # A tibble: 336,776 x 6  
##   year month   day hour minute departure  
##   <int> <int> <int> <dbl> <dbl>      <dtm>  
## 1  2013     1     1     5     15 2013-01-01 05:15:00  
## 2  2013     1     1     5     29 2013-01-01 05:29:00  
## 3  2013     1     1     5     40 2013-01-01 05:40:00  
## 4  2013     1     1     5     45 2013-01-01 05:45:00  
## 5  2013     1     1     6      0 2013-01-01 06:00:00  
## 6  2013     1     1     5     58 2013-01-01 05:58:00  
## 7  2013     1     1     6      0 2013-01-01 06:00:00  
## 8  2013     1     1     6      0 2013-01-01 06:00:00  
## 9  2013     1     1     6      0 2013-01-01 06:00:00
```

Going in reverse

```
datetime <- ymd_hms("2016-07-08 12:34:56")
```

```
year(datetime)
```

```
## [1] 2016
```

```
month(datetime)
```

```
## [1] 7
```

```
mday(datetime)
```

```
## [1] 8
```

Calculating Time Spans

- Really common situation for me: Dataset like the below, need to calculate number of weeks/months, etc., either between dates, or from a specific date.

```
sid <- rep(1:4, each = 3)
date <- c("9/3/08", "12/10/08", "4/22/09", "8/29/08", "12/5/08", "4/17/09", "8/29/08", "12/4/08", "4/23/09", "9/3/08", "12/1/08")
score <- c(222, 225, 223, 194, 196, 201, 194, 209, 197, 191, 197, 214)
d <- data.frame(sid = sid, date = date, score = score)
d
```

##	sid	date	score
## 1	1	9/3/08	222
## 2	1	12/10/08	225
## 3	1	4/22/09	223
## 4	2	8/29/08	194
## 5	2	12/5/08	196
## 6	2	4/17/09	201
## 7	3	8/29/08	194
## 8	3	12/4/08	209
## 9	3	4/23/09	197
## 10	4	9/3/08	191
## 11	4	12/1/08	197

First - convert to date

```
d <- d %>%  
  mutate(date = mdy(date))  
d
```

##	sid	date	score
## 1	1	2008-09-03	222
## 2	1	2008-12-10	225
## 3	1	2009-04-22	223
## 4	2	2008-08-29	194
## 5	2	2008-12-05	196
## 6	2	2009-04-17	201
## 7	3	2008-08-29	194
## 8	3	2008-12-04	209
## 9	3	2009-04-23	197
## 10	4	2008-09-03	191
## 11	4	2008-12-01	197
## 12	4	2009-04-20	214

What to compute from?

- In my case, I often want to calculate the date from the first day of the school year.
- First, create a date object with that date

```
first_day <- mdy("08/05/2008")  
first_day
```

```
## [1] "2008-08-05"
```

- Next, compute the difference between that date and the corresponding date the test was administered.

```
d %>%  
  mutate(days_elapsed = date - first_day)
```

##	sid	date	score	days_elapsed
## 1	1	2008-09-03	222	29 days
## 2	1	2008-12-10	225	127 days
## 3	1	2009-04-22	223	260 days
## 4	2	2008-08-29	194	24 days
## 5	2	2008-12-05	196	122 days
## 6	2	2009-04-17	201	255 days
## 7	3	2008-08-29	194	24 days
## 8	3	2008-12-04	209	121 days
## 9	3	2009-04-23	197	261 days
## 10	4	2008-09-03	191	29 days
## 11	4	2008-12-01	197	118 days
## 12	4	2009-04-20	214	258 days

Take a second...

What if I wanted to calculate date from the first assessment?

One method

```
d %>%
  group_by(sid) %>%
  arrange(date) %>%
  mutate(first_date = first(date),
         days_elapsed = date - first_date) %>%
  arrange(sid)
```

```
## # A tibble: 12 x 5
## # Groups:   sid [4]
##      sid      date score first_date days_elapsed
##   <int>   <date> <dbl>     <date>      <time>
## 1     1 2008-09-03   222 2008-09-03        0 days
## 2     1 2008-12-10   225 2008-09-03       98 days
## 3     1 2009-04-22   223 2008-09-03      231 days
## 4     2 2008-08-29   194 2008-08-29        0 days
## 5     2 2008-12-05   196 2008-08-29       98 days
## 6     2 2009-04-17   201 2008-08-29      231 days
## 7     3 2008-08-29   194 2008-08-29        0 days
## 8     3 2008-12-04   209 2008-08-29       97 days
## 9     3 2009-04-23   197 2008-08-29      237 days
## 10    4 2008-09-03   191 2008-09-03        0 days
```

What if I wanted days between each assessment?

- Some knowledge of base R comes in handy here: `lag`

```
d %>%  
  group_by(sid) %>%  
  arrange(date) %>%  
  mutate(days_between = date - lag(date)) %>%  
  arrange(sid)
```

```
## # A tibble: 12 x 4  
## # Groups:   sid [4]  
##      sid      date score days_between  
##   <int>   <date> <dbl>      <time>  
## 1     1 2008-09-03   222      NA days  
## 2     1 2008-12-10   225      98 days  
## 3     1 2009-04-22   223     133 days  
## 4     2 2008-08-29   194      NA days  
## 5     2 2008-12-05   196      98 days  
## 6     2 2009-04-17   201     133 days  
## 7     3 2008-08-29   194      NA days  
## 8     3 2008-12-04   209      97 days  
## 9     3 2009-04-23   197     140 days
```

Different metric?

- Suppose I instead wanted weeks

```
first_day_weeks <- week(first_day)
first_day_weeks
```

```
## [1] 32
```

```
d <- d %>%
  mutate(weeks_elapsed = week(date) - first_day_weeks)
d
```

```
## Source: local data frame [12 x 5]
## Groups: sid [4]
##
## # A tibble: 12 x 5
##       sid      date score occasion weeks_elapsed
##   <int>   <date> <dbl>     <int>         <dbl>
## 1     1 2008-09-03   222         1           4
## 2     1 2008-12-10   225         2          18
## 3     1 2009-04-22   223         3          -16
## 4     2 2008-08-29   194         1           3
```

Check

Uh oh...What to do?

```
d %>%
  mutate(days_elapsed = date - first_day,
         check = days_elapsed / 7) %>%
  select(weeks_elapsed, check)
```

```
## Source: local data frame [12 x 3]
## Groups: sid [4]
##
## # A tibble: 12 x 3
##       sid weeks_elapsed      check
##   <int>      <dbl>      <time>
## 1     1         4 4.142857 days
## 2     1        18 18.142857 days
## 3     1       -16 37.142857 days
## 4     2         3 3.428571 days
## 5     2        17 17.428571 days
## 6     2       -16 36.428571 days
## 7     3         3 3.428571 days
## 8     3        17 17.285714 days
## 9     3       -15 37.285714 days
```

What about months?

One method...

```
first_day_months <- month(first_day)
first_day_months
```

```
## [1] 8
```

```
d <- d %>%
  mutate(months_elapsed = ifelse(year(date) == "2008",
                                month(date) - first_day_months,
                                (month(date) - first_day_months) + 12))
```

d

```
## Source: local data frame [12 x 6]
## Groups: sid [4]
##
## # A tibble: 12 x 6
##       sid      date score occasion weeks_elapsed months_elapsed
##   <int>   <date> <dbl>    <int>      <dbl>        <dbl>
## 1     1 2008-09-03   222        1         4          1
## 2     1 2008-12-10   225        2        18          4
## 3     1 2009-04-22   223        3       -16          8
## 4     2 2008-08-29   194        1         3          0
## 5     2 2008-12-05   196        2        17          4
## 6     2 2009-04-17   201        3       -16          8
## 7     3 2008-08-29   194        1         3          0
## 8     3 2008-12-04   209        2        17          4
## 9     3 2009-04-23   197        3       -15          8
## 10    4 2008-09-03   191        1         4          1
## 11    4 2008-12-01   197        2        16          4
## 12    4 2009-04-20   214        3       -16          8
```

Alternative

Non-tidyverse package but useful for months, specifically *mondate*

```
# install.packages("mondate")  
library(mondate)  
first_day_mondate <- as.mondate(first_day)  
first_day_mondate
```

```
## mondate: timeunits="months"  
## [1] 08/05/2008
```

```
d <- d %>%
  mutate(mondate = as.mondate(date),
         months_elapsed2 = mondate - first_day_mondate)
d
```

##	sid	date	score	months_elapsed	mondate	months_elapsed2
## 1	1	2008-09-03	222	1	09/03/2008	0.9387097 months
## 2	1	2008-12-10	225	4	12/10/2008	4.1612903 months
## 3	1	2009-04-22	223	8	04/22/2009	8.5720430 months
## 4	2	2008-08-29	194	0	08/29/2008	0.7741935 months
## 5	2	2008-12-05	196	4	12/05/2008	4.0000000 months
## 6	2	2009-04-17	201	8	04/17/2009	8.4053763 months
## 7	3	2008-08-29	194	0	08/29/2008	0.7741935 months
## 8	3	2008-12-04	209	4	12/04/2008	3.9677419 months
## 9	3	2009-04-23	197	8	04/23/2009	8.6053763 months
## 10	4	2008-09-03	191	1	09/03/2008	0.9387097 months
## 11	4	2008-12-01	197	4	12/01/2008	3.8709677 months
## 12	4	2009-04-20	214	8	04/20/2009	8.5053763 months

A few last notes on dates

- *lubridate* provides **duration** and **period** classes that may be helpful
 - durations are always reported in seconds
- Periods help account for things like time zones and leap years

Use durations to calculate dates

```
today() + ddays(123)
```

```
## [1] "2017-09-29"
```

```
today() + dweeks(1)
```

```
## [1] "2017-06-05"
```

```
today() - dyears(1)
```

```
## [1] "2016-05-29"
```

Another alternative for months

```
d %>%  
  mutate(days_elapsed = date - first_day,  
         seconds_elapsed = as.duration(days_elapsed),  
         months_elapsed3 = seconds_elapsed / 2.628e+6) %>%  
  select(contains("month"))
```

```
##      months_elapsed  months_elapsed2    months_elapsed3  
## 1                1 0.9387097 months 0.953424657534247s  
## 2                4 4.1612903 months 4.17534246575342s  
## 3                8 8.5720430 months 8.54794520547945s  
## 4                0 0.7741935 months 0.789041095890411s  
## 5                4 4.0000000 months 4.01095890410959s  
## 6                8 8.4053763 months 8.38356164383562s  
## 7                0 0.7741935 months 0.789041095890411s  
## 8                4 3.9677419 months 3.97808219178082s  
## 9                8 8.6053763 months 8.58082191780822s  
## 10               1 0.9387097 months 0.953424657534247s  
## 11               4 3.8709677 months 3.87945205479452s  
## 12               8 8.5053763 months 8.48219178082192s
```

periods

```
one_pm <- ymd_hms("2016-03-12 13:00:00", tz = "America/New_York")  
one_pm
```

```
## [1] "2016-03-12 13:00:00 EST"
```

```
one_pm + ddays(1)
```

```
## [1] "2016-03-13 14:00:00 EDT"
```

```
one_pm + days(1)
```

```
## [1] "2016-03-13 13:00:00 EDT"
```

Summary

- Dates are harder than expected
 - time zones
 - leap years
 - daylight savings, etc.
- *lubridate* can help, but you always need to be careful
- We didn't talk about calculating seconds, milliseconds, etc., but that's easily done as well.